

# Bases de Dados



Universidade do Porto  
Faculdade de Engenharia  
**FEUP**

1

## NoSQL

Observação: baseado em <http://www.stanford.edu/class/cs145/ppt/cs145nosql.pptx>

# SQL Characteristics



- Data stored in columns and tables
- Relationships represented by data
- Data Manipulation Language
  - `Select T1.Column1, T2.Column2 ...`  
`From Table1, Table2 ...`  
`Where T1.Column1 = T2.Column1 ...`
- Data Definition Language
  - Schema defined at the start
  - `Create Table (Column1 Datatype1, Column2 Datatype 2, ...)`
  - Constraints to define and enforce relationships
    - ✦ Primary Key
    - ✦ Foreign Key
    - ✦ Etc.
  - Triggers to respond to Insert, Update , & Delete
  - Stored Modules
  - `Alter ...`
  - `Drop ...`
  - Security and Access Control
- Transactions: ACID properties
- Abstraction from physical layer

# NoSQL Definition



From [www.nosql-database.org](http://www.nosql-database.org):

Next Generation Databases mostly addressing some of the points: being **non-relational, distributed, open-source** and **horizontal scalable**. The original intention has been **modern web-scale databases**. The movement began early 2009 and is growing rapidly. Often more characteristics apply as: **schema-free, easy replication support, simple API, eventually consistent / BASE** (not ACID), a **huge data amount**, and more.

# Why NoSQL?

4

- Not every data management/analysis problem is best solved using a traditional relational DBMS
  - e.g., remember text search?
- “NoSQL” = “Not only SQL”
  - ... = Not using traditional relational DBMS

# NoSQL Systems

5

- Alternative to traditional relational DBMS
  - + Flexible schema
    - ✦ including unstructured documents, images, videos
  - + Quicker/cheaper to set up
  - + Massive scalability
  - + Relaxed consistency
    - ✦ higher performance & availability but fewer guarantees
  - No declarative query language
    - ✦ more programming

# Example of Advanced Query (1/2)

6

- Social-network graph
  - Each record: UserID1, UserID2
  - Separate records: UserID, name, age, gender, ...
- Task
  - Find all friends of friends of friends of ... friends of given user

# Example of Advanced Query (2/2)

7

- Wikipedia pages
  - Large collection of documents
  - Combination of structured and unstructured data
- Task
  - Retrieve introductory paragraph of all pages about U.S. presidents before 1900

# Example in MongoDB

8

- find all orders shipped by CTT-Expresso  
`db.order.find( { shipping: {carrier: "ctt-expresso"} } )`
- ... and process results  

```
var cursor = db.order.find( { shipping: {carrier: "ctt-expresso"} } );
cursor.hasNext();
cursor.forEach(
    function(item) {
        print(tojson(item))
    });
```



# Types of NoSQL Systems

9

- MapReduce framework
- Key-value stores
- Document stores
- Graph database systems

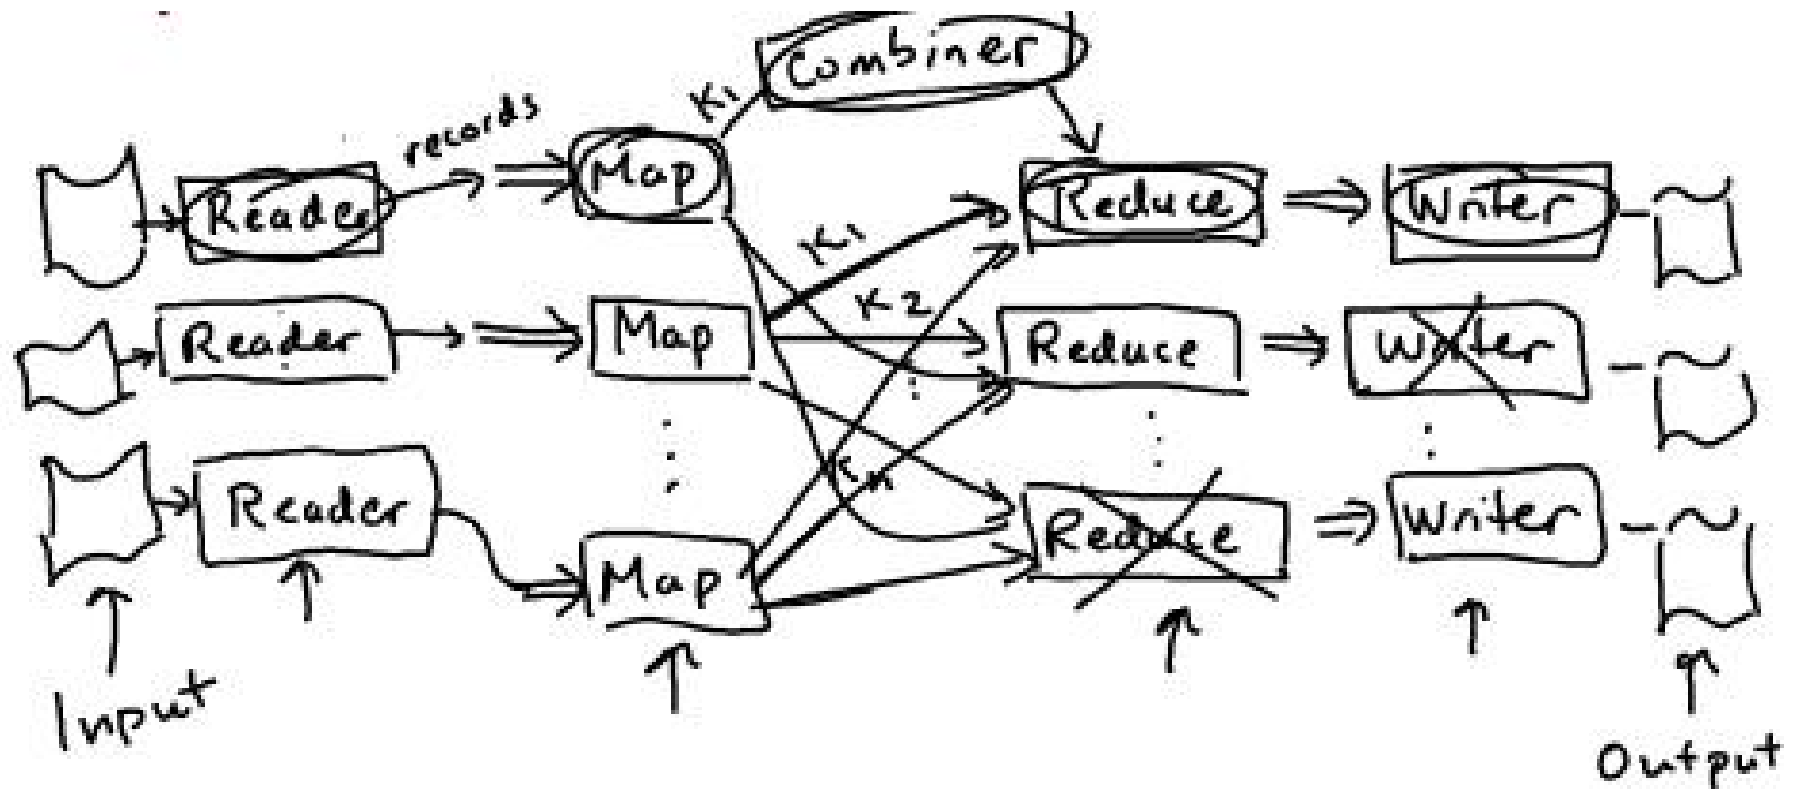
# MapReduce

10

- Originally from Google, open source Hadoop
  - No data model, data stored in files
  - User provides specific functions
    - ✦ map()
    - ✦ reduce()
- System provides
  - data processing “glue”
  - fault-tolerance
  - scalability

# MapReduce Architecture

11



# MapReduce Framework

12

- Schemas and declarative queries are missed
  - Hive
    - ✦ schemas
    - ✦ SQL-like query language
  - Pig
    - ✦ more imperative but with relational operators
- Both compile to “workflow” of Hadoop (MapReduce) jobs

# Key-Value Stores

13

- Extremely simple interface
  - Data model: (key, value) pairs
  - Operations:
    - ✦ Insert(key,value)
    - ✦ Fetch(key)
- Implementation
  - efficiency
  - scalability
  - fault-tolerance
- Example systems
  - Google BigTable
  - Amazon Dynamo
  - Cassandra
  - Voldemort
  - HBase

# Document Stores

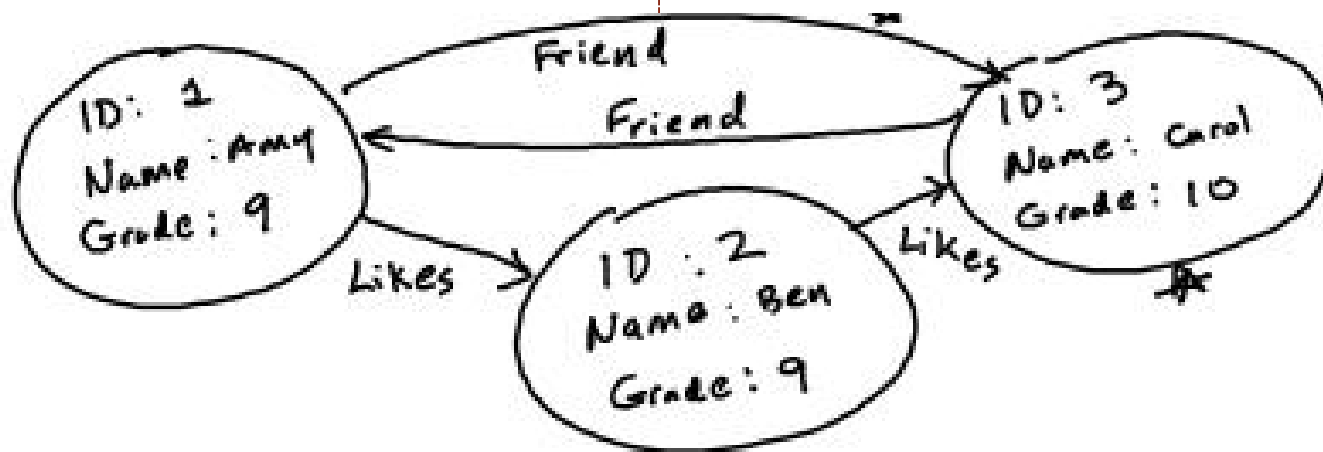
14

- Key-Value Stores where value is a document
    - Data model: (key, document) pairs
    - Document: JSON, XML, other semistructured formats
  - Basic operations:
    - Insert(key,document)
    - Fetch(key)
  - Also Fetch based on document contents
- Example systems
    - CouchDB
    - MongoDB
    - SimpleDB

# Graph Database Systems

15

- Data model: nodes and edges
  - Nodes may have properties (including ID)
  - Edges may have labels or roles
- Example systems
  - Neo4j
  - FlockDB
  - Pregel



# Distributed systems



A **distributed system** should have the following characteristics:

- **Consistency**
  - All nodes see the same data at the same time – Wikipedia
  - Client perceives that a set of operations has occurred all at once – Pritchett
  - More like Atomic in ACID transaction properties
- **Availability**
  - Node failures do not prevent survivors from continuing to operate – Wikipedia
  - Every operation must terminate in an intended response – Pritchett
- **Partition tolerance**
  - The system continues to operate despite arbitrary message loss – Wikipedia
  - Operations will complete, even if individual components are unavailable – Pritchett

<http://queue.acm.org/detail.cfm?id=1394128>



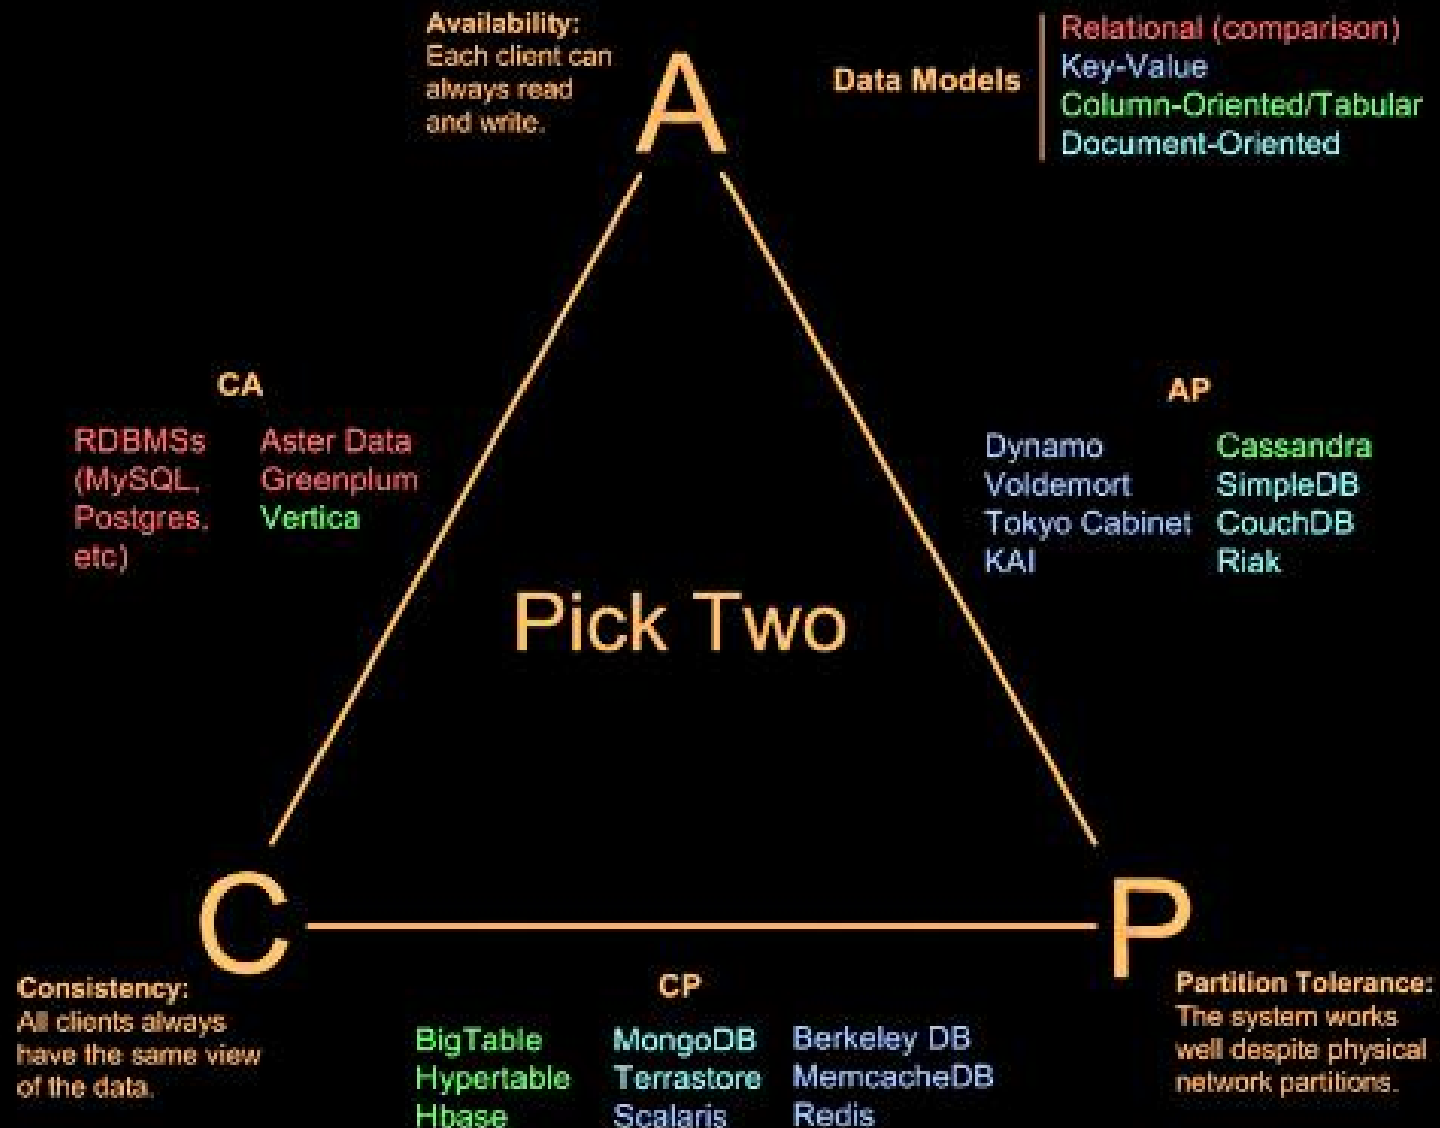
# Brewer's CAP Theorem



A **distributed system** can support only two of the following characteristics:

- Consistency
- Availability
- Partition tolerance

# Visual Guide to NoSQL Systems



# Summary

19

- “NoSQL” = “Not Only SQL”
  - Not every data management/analysis problem is best solved exclusively using a traditional DBMS
- Current incarnations
  - MapReduce framework
  - Key-value stores
  - Document stores
  - Graph database systems

